

In Previous update we did transfer learning using UNET which is trained on imagenet Dataset. In mininetwork we used same UNET architecture(ecoder, decoder) from scrath to create a mininetwork.

```
In [ ]: !nvidia-smi
```

```
Tue May  9 16:44:50 2023

+-----+
| NVIDIA-SMI 525.85.12      Driver Version: 525.85.12      CUDA Version: 12.0      |
+-----+-----+-----+-----+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M.         |
+-----+-----+-----+-----+-----+-----+
|  0  Tesla T4               Off        | 00000000:00:04.0 Off |                    0 |
| N/A   51C    P8             10W /  70W |  0MiB / 15360MiB |      0%      Default |
+-----+-----+-----+-----+-----+-----+

+-----+
| Processes:                                                       GPU Memory |
|  GPU   GI    CI          PID    Type   Process name                  Usage      |
|-----+-----+-----+-----+-----+-----+
| No running processes found                                     |
+-----+
```

```
In [ ]: import numpy as np
import pandas as pd
import os, glob
from torch.utils.data import Dataset
import torch
from PIL import Image
import matplotlib.pyplot as plt
from albumentations.pytorch import ToTensorV2
import albumentations as A
```

```
In [ ]: import torch.nn as nn
from torch.optim import Adam
from tqdm import tqdm
```

```
In [ ]: IMAGE_DIR = "/content/drive/MyDrive/Unet/Dataset_AL/images"
MASKS_DIR = "/content/drive/MyDrive/Unet/Dataset_AL/annotations"
```

```
In [ ]: IMAGE_PATHS = glob.glob(IMAGE_DIR + "/*")
MASKS_PATHS = glob.glob(MASKS_DIR + "/*")
```

```
IMAGE_PATHS.sort()
MASKS_PATHS.sort()
```

```
In [ ]: IMAGE_PATHS[10], MASKS_PATHS[10]
```

```
Out[ ]: ('/content/drive/MyDrive/Unet/Dataset_AL/images/UNH_DL_100.jpg',
        '/content/drive/MyDrive/Unet/Dataset_AL/annotations/UNH_DL_100.png')
```

```
In [ ]: df = pd.DataFrame({
        'image_path' : IMAGE_PATHS,
        'mask_path'  : MASKS_PATHS
    })

df.head()
```

```
Out[ ]:
```

	image_path	mask_path
0	/content/drive/MyDrive/Unet/Dataset_AL/images/...	/content/drive/MyDrive/Unet/Dataset_AL/annotat...
1	/content/drive/MyDrive/Unet/Dataset_AL/images/...	/content/drive/MyDrive/Unet/Dataset_AL/annotat...
2	/content/drive/MyDrive/Unet/Dataset_AL/images/...	/content/drive/MyDrive/Unet/Dataset_AL/annotat...
3	/content/drive/MyDrive/Unet/Dataset_AL/images/...	/content/drive/MyDrive/Unet/Dataset_AL/annotat...
4	/content/drive/MyDrive/Unet/Dataset_AL/images/...	/content/drive/MyDrive/Unet/Dataset_AL/annotat...

```
In [ ]: df.shape
```

```
Out[ ]: (130, 2)
```

Dataset

```
In [ ]: class UnetDataset(Dataset):
        def __init__(self, df, transform = None):

            self.df = df
            self.transforms = transform

        def __len__(self):
            return len(self.df)

        def __getitem__(self, idx):

            img = np.array(Image.open(self.df.iloc[idx]['image_path']))
            mask = np.array(Image.open(self.df.iloc[idx]['mask_path']))

            if self.transforms is not None:
                aug = self.transforms(image=img, mask=mask)
                img = aug['image']
                mask = aug['mask']
                # mask = torch.max(mask, dim=2)[0]
                mask = mask.long()
```

```
return img, mask
```

Image mean and std

```
In [ ]: transforms_tmp = A.Compose([
        A.Resize(width=256, height=256),
        A.Normalize(mean = (0, 0, 0), std = (1, 1, 1)),
        ToTensorV2(),
    ])
```

```
In [ ]: d = UnetDataset(df, transform=transforms_tmp)
        dataloader = torch.utils.data.DataLoader(d,
                                                    batch_size=1,
                                                    shuffle=True,
                                                    num_workers=2)
```

```
In [ ]: # placeholders
        psum      = torch.tensor([0.0, 0.0, 0.0])
        psum_sq   = torch.tensor([0.0, 0.0, 0.0])

        for i, (image, target) in enumerate(tqdm(dataloader)):

            psum      += image.sum(axis      = [0, 2, 3])
            psum_sq   += (image ** 2).sum(axis = [0, 2, 3])
```

```
100%|██████████| 130/130 [00:16<00:00, 7.93it/s]
```

```
In [ ]: ##### FINAL CALCULATIONS

        # pixel count
        count = len(df) * 256 * 256

        # mean and std
        total_mean = psum / count
        total_var  = (psum_sq / count) - (total_mean ** 2)
        total_std  = torch.sqrt(total_var)

        # output
        print('mean: ' + str(total_mean))
        print('std: ' + str(total_std))
```

```
mean: tensor([0.4893, 0.4996, 0.4967])
std:  tensor([0.2438, 0.2386, 0.2563])
```

Transforms

```
In [ ]: import albumentations as A
        from albumentations.pytorch import ToTensorV2

        # Define a list of augmentations to apply
        transforms_train = A.Compose([
            A.Resize(width=256, height=256),
```

```

        A.HorizontalFlip(p=0.5),
        A.Rotate(limit=30, p=0.5),
        A.RandomBrightnessContrast(p=0.2),
        A.Blur(p=0.1),
        A.Normalize(total_mean, total_std),
        ToTensorV2(),
    ])

transforms_val = A.Compose([
    A.Resize(width=256, height=256),
    A.Normalize(total_mean, total_std),
    ToTensorV2(),
])

```

Plot

```

In [ ]: d = UnetDataset(df, transform=transforms_train)
        image, mask = d.__getitem__(15)

        # Convert the tensors back to NumPy arrays
        image = image.permute(1, 2, 0).numpy()
        mask = mask.numpy()

        # Plot the image and mask side-by-side
        fig, ax = plt.subplots(1, 2, figsize=(10, 5))

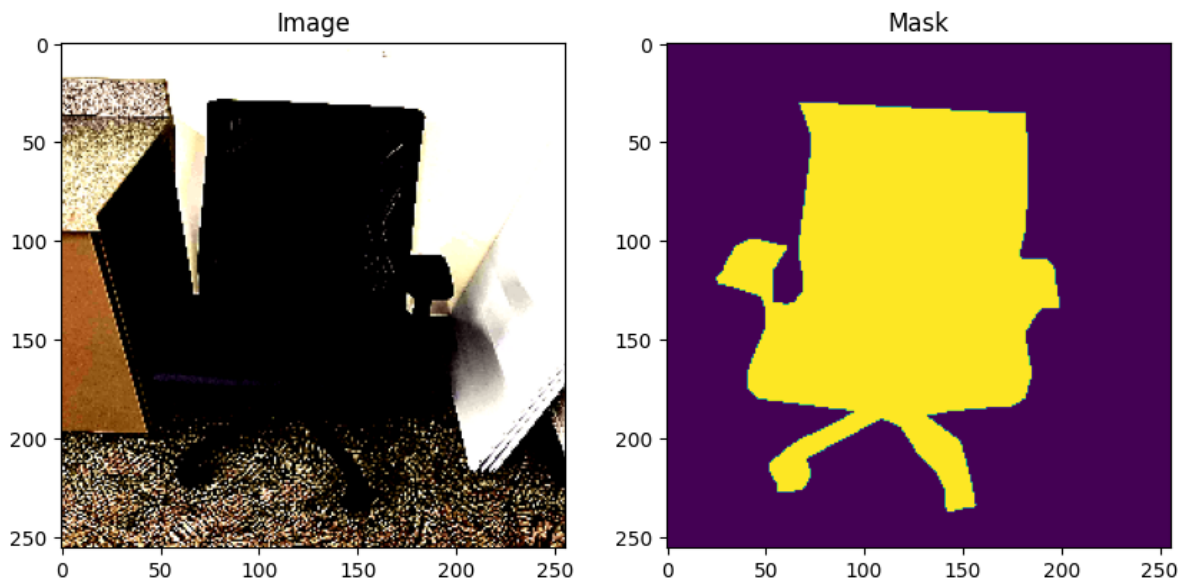
        ax[0].imshow(image)
        ax[0].set_title("Image")

        ax[1].imshow(mask)
        ax[1].set_title("Mask")

        plt.show()

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Train Valid Test split

```
In [ ]: # Lengths of training, validation, and test sets
data_length = len(df)
print(data_length)
train_length = int(0.8 * data_length)
validation_length = int(0.1 * data_length)
test_length = data_length - train_length - validation_length

130
```

```
In [ ]: train_length, validation_length, test_length
```

```
Out[ ]: (104, 13, 13)
```

```
In [ ]: from sklearn.model_selection import train_test_split

train_df, valid_df = train_test_split(df, test_size=0.1, shuffle=True)
train_df, test_df = train_test_split(train_df, test_size=0.11, shuffle=True)

train_df.shape, valid_df.shape, test_df.shape
```

```
Out[ ]: ((104, 2), (13, 2), (13, 2))
```

Model

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
In [ ]: class encoding_block(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(encoding_block, self).__init__()
        model = []
```

```

        model.append(nn.Conv2d(in_channels, out_channels, 3, 1, 1, bias=False))
        model.append(nn.BatchNorm2d(out_channels))
        model.append(nn.ReLU(inplace=True))
        model.append(nn.Conv2d(out_channels, out_channels, 3, 1, 1, bias=False))
        model.append(nn.BatchNorm2d(out_channels))
        model.append(nn.ReLU(inplace=True))
        self.conv = nn.Sequential(*model)
    def forward(self, x):
        return self.conv(x)

```

```

In [ ]: class UNet(nn.Module):
    def __init__(self, out_channels=2, features=[64, 128, 256, 512]):
        super(UNet, self).__init__()
        self.pool = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2))
        self.conv1 = encoding_block(3, features[0])
        self.conv2 = encoding_block(features[0], features[1])
        self.conv3 = encoding_block(features[1], features[2])
        self.conv4 = encoding_block(features[2], features[3])
        self.conv5 = encoding_block(features[3]*2, features[3])
        self.conv6 = encoding_block(features[3], features[2])
        self.conv7 = encoding_block(features[2], features[1])
        self.conv8 = encoding_block(features[1], features[0])
        self.tconv1 = nn.ConvTranspose2d(features[-1]*2, features[-1], kernel_size=
self.tconv2 = nn.ConvTranspose2d(features[-1], features[-2], kernel_size=2,
self.tconv3 = nn.ConvTranspose2d(features[-2], features[-3], kernel_size=2,
self.tconv4 = nn.ConvTranspose2d(features[-3], features[-4], kernel_size=2,
self.bottleneck = encoding_block(features[3], features[3]*2)
self.final_layer = nn.Conv2d(features[0], out_channels, kernel_size=1)

    # Initialize the weights and biases
    self._initialize_weights()

    def _initialize_weights(self):
        for m in self.modules():
            if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
                nn.init.kaiming_normal_(m.weight, mode='fan_in', nonlinearity='relu')
                if m.bias is not None:
                    nn.init.constant_(m.bias, 0)
            elif isinstance(m, nn.BatchNorm2d):
                nn.init.constant_(m.weight, 1)
                nn.init.constant_(m.bias, 0)

    def forward(self, x):
        skip_connections = []
        x = self.conv1(x)
        skip_connections.append(x)
        x = self.pool(x)
        x = self.conv2(x)
        skip_connections.append(x)
        x = self.pool(x)
        x = self.conv3(x)
        skip_connections.append(x)
        x = self.pool(x)
        x = self.conv4(x)
        skip_connections.append(x)

```

```

x = self.pool(x)
x = self.bottleneck(x)
skip_connections = skip_connections[:-1]
x = self.tconv1(x)
x = torch.cat((skip_connections[0], x), dim=1)
x = self.conv5(x)
x = self.tconv2(x)
x = torch.cat((skip_connections[1], x), dim=1)
x = self.conv6(x)
x = self.tconv3(x)
x = torch.cat((skip_connections[2], x), dim=1)
x = self.conv7(x)
x = self.tconv4(x)
x = torch.cat((skip_connections[3], x), dim=1)
x = self.conv8(x)
x = self.final_layer(x)
return x

```

```

In [ ]: model = UNet(out_channels=5+1)

input = torch.rand(1, 3, 256, 256)
output = model(input)

```

```

In [ ]: input.shape, output.shape

```

```

Out[ ]: (torch.Size([1, 3, 256, 256]), torch.Size([1, 6, 256, 256]))

```

Train Valid Test Functions

```

In [ ]: def train_epoch(model, optimizer, criterion, data_loader, device):
    # Set the model to training mode
    model.train()

    # Initialize variables to keep track of loss and accuracy
    running_loss = 0
    num_correct = 0
    num_pixels = 0
    dice_score = 0

    # Loop over the data loader
    for batch_idx, (inputs, targets) in enumerate(tqdm(data_loader)):
        # Move the inputs and targets to the device
        inputs, targets = inputs.to(device), targets.to(device)

        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, targets)

        # Backward pass
        loss.backward()
        optimizer.step()

```

```

# Update the running loss and accuracy
running_loss += loss.item() * inputs.size(0)

softmax = nn.Softmax(dim=1)
preds = torch.argmax(softmax(outputs),axis=1)
num_correct += (preds == targets).sum()
num_pixels += torch.numel(preds)
dice_score += (2 * (preds * targets).sum()) / ((preds + targets).sum() + 1e-6)

# Calculate the average loss and accuracy for the epoch
avg_loss = running_loss / len(data_loader.dataset)
dice_score = dice_score / len(data_loader.dataset)

print("Train loss : ", avg_loss)
print("Train dice score : ", dice_score.item())

return avg_loss, dice_score.item()

```

```

In [ ]: def valid_epoch(model, optimizer, criterion, data_loader, device):
    # Set the model to training mode
    model.eval()

    # Initialize variables to keep track of loss and accuracy
    running_loss = 0
    num_correct = 0
    num_pixels = 0
    dice_score = 0

    # Disable gradient calculation to speed up inference
    with torch.no_grad():
        # Loop over the data loader
        for batch_idx, (inputs, targets) in enumerate(tqdm(data_loader)):
            # Move the inputs and targets to the device
            inputs, targets = inputs.to(device), targets.to(device)

            # Zero the gradients
            optimizer.zero_grad()

            # Forward pass
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            # Update the running loss and accuracy
            running_loss += loss.item() * inputs.size(0)

            softmax = nn.Softmax(dim=1)
            preds = torch.argmax(softmax(outputs),axis=1)
            num_correct += (preds == targets).sum()
            num_pixels += torch.numel(preds)
            dice_score += (2 * (preds * targets).sum()) / ((preds + targets).sum() + 1e-6)

    # Calculate the average loss and accuracy for the epoch
    avg_loss = running_loss / len(data_loader.dataset)

```



```

dice_score = dice_score / len(data_loader.dataset)

print("Valid loss : ", avg_loss)
print("Valid dice score : ", dice_score.item())

return avg_loss, dice_score.item()

```

Data Loaders

```

In [ ]: DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
LEARNING_RATE = 1e-3
WEIGHT_DECAY = 5e-4
BATCH_SIZE = 4
NUM_WORKERS = 4
CHECKPOINT_FILE = "Best.pth.tar"
PIN_MEMORY = True
SAVE_MODEL = True
LOAD_MODEL = False

```

```

In [ ]: from torch.utils.data import Dataset, DataLoader

```

```

In [ ]: train_dataset = UnetDataset(train_df, transform=transforms_train)
train_loader = DataLoader(dataset=train_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=True,
                           num_workers=2,
                           pin_memory=PIN_MEMORY)

valid_dataset = UnetDataset(valid_df, transform=transforms_val)
valid_loader = DataLoader(dataset=valid_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=False,
                           pin_memory=PIN_MEMORY)

test_dataset = UnetDataset(test_df, transform=transforms_val)
test_loader = DataLoader(dataset=test_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=False,
                           pin_memory=PIN_MEMORY)

```

```

In [ ]: def save_checkpoint(state, filename="model.pth.tar"):
    print("=> Saving checkpoint")
    torch.save(state, filename)

```

```

In [ ]: def load_checkpoint(checkpoint, model):
    print("=> Loading checkpoint")
    model.load_state_dict(checkpoint["state_dict"])

```

Model and optimizer and loss

```

In [ ]: # model

```

```

model = UNet(out_channels=5+1)
model.to(DEVICE)

optimizer = torch.optim.AdamW(params=model.parameters(),
                               lr=LEARNING_RATE,
                               weight_decay=WEIGHT_DECAY)

criterion = nn.CrossEntropyLoss()

```

Overfitting model on training data

```

In [ ]: train_dataset1 = UnetDataset(train_df.sample(1), transform=transforms_train)
train_loader1 = DataLoader(dataset=train_dataset1,
                           batch_size=BATCH_SIZE,
                           shuffle=True,
                           num_workers=2,
                           pin_memory=PIN_MEMORY)

for epoch in range(10):
    print("Epoch = ", epoch)

    train_loss, train_score = train_epoch(
        model,
        optimizer=optimizer,
        criterion=criterion,
        data_loader=train_loader1,
        device=DEVICE
    )

```

Epoch = 0

100%|██████████| 1/1 [00:00<00:00, 1.74it/s]

Train loss : 1.6151604652404785

Train dice score : 0.47445040941238403

Epoch = 1

100%|██████████| 1/1 [00:00<00:00, 1.82it/s]

Train loss : 1.4422485828399658

Train dice score : 0.6047606468200684

Epoch = 2

100%|██████████| 1/1 [00:00<00:00, 1.82it/s]

Train loss : 1.1793590784072876

Train dice score : 0.8311318755149841

Epoch = 3

100%|██████████| 1/1 [00:00<00:00, 1.85it/s]

Train loss : 1.0406230688095093

Train dice score : 0.8936702013015747

Epoch = 4

100%|██████████| 1/1 [00:00<00:00, 1.80it/s]

Train loss : 0.9175289869308472

Train dice score : 0.8488644361495972

Epoch = 5

100%|██████████| 1/1 [00:00<00:00, 1.85it/s]

Train loss : 0.8025119304656982

Train dice score : 0.8418070077896118

Epoch = 6

100%|██████████| 1/1 [00:00<00:00, 1.85it/s]

Train loss : 0.7840336561203003

Train dice score : 0.8662620782852173

Epoch = 7

100%|██████████| 1/1 [00:00<00:00, 1.88it/s]

Train loss : 0.751652717590332

Train dice score : 0.899687647819519

Epoch = 8

100%|██████████| 1/1 [00:00<00:00, 1.89it/s]

Train loss : 0.7180434465408325

Train dice score : 0.9020321369171143

Epoch = 9

100%|██████████| 1/1 [00:00<00:00, 1.88it/s]

Train loss : 0.6745339632034302

Train dice score : 0.8812836408615112

```
In [ ]: softmax = nn.Softmax(dim=1)
# Disable gradient calculation to speed up inference
with torch.no_grad():
    # Loop over the data loader
    for batch_idx, (inputs, targets) in enumerate(tqdm(train_loader1)):
        # Move the inputs and targets to the device
        inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)

        # Forward pass
        outputs = model(inputs)
        outputs = torch.argmax(softmax(outputs), axis=1)

        for i in range(inputs.shape[0]):

            image = inputs[i].to('cpu')
            mask = targets[i].to('cpu')
            pred_mask = outputs[i].to('cpu')

            fig, ax = plt.subplots(1, 3, figsize=(18, 18))
            softmax = nn.Softmax(dim=1)

            image = image.permute(1, 2, 0).numpy()
            mask = mask.numpy()
            pred_mask = pred_mask.numpy()

            ax[0].imshow(image)
            ax[0].set_title("Image")


            ax[1].imshow(mask)
            ax[1].set_title("Mask")

            ax[2].imshow(pred_mask)
            ax[2].set_title("Predicted Mask")
```

```
plt.show()
```

```
# break
```

```
0%|          | 0/1 [00:00<?, ?it/s]WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
```



```
100%|██████████| 1/1 [00:01<00:00, 1.40s/it]
```

Hyperparameter tuning

learning rate, momentum, and weight decay

```
In [ ]: best_score = -np.inf

best_lr = None
best_weight_decay = None

for lr in [ 0.001, 0.01, 0.1]:
    for weight_decay in [0.0001, 0.001, 0.01]:

        model = UNet(out_channels=5+1)
        model.to(DEVICE)

        optimizer = torch.optim.AdamW(params=model.parameters(),
                                       lr=lr,
                                       weight_decay=weight_decay)

        criterion = nn.CrossEntropyLoss()

        train_loss, train_score = train_epoch(
            model,
            optimizer=optimizer,
            criterion=criterion,
            data_loader=train_loader,
            device=DEVICE
        )

        valid_loss, valid_score = valid_epoch(
            model,
```

```

optimizer=optimizer,
criterion=criterion,
data_loader=valid_loader,
device=DEVICE
)

if valid_score > best_score:
    best_score = valid_score

    best_lr = lr
    best_weight_decay = weight_decay

```

```
best_lr, best_weight_decay
```

```
100%|██████████| 26/26 [00:15<00:00, 1.72it/s]
```

```
Train loss : 1.2057688465485206
```

```
Train dice score : 0.1017453745007515
```

```
100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
```

```
Valid loss : 3.228561364687406
```

```
Valid dice score : 0.1436549574136734
```

```
100%|██████████| 26/26 [00:15<00:00, 1.72it/s]
```

```
Train loss : 1.3119862492267902
```

```
Train dice score : 0.0923336073756218
```

```
100%|██████████| 4/4 [00:03<00:00, 1.30it/s]
```

```
Valid loss : 5.940927844781142
```

```
Valid dice score : 0.16841505467891693
```

```
100%|██████████| 26/26 [00:14<00:00, 1.79it/s]
```

```
Train loss : 1.229852781845973
```

```
Train dice score : 0.12860193848609924
```

```
100%|██████████| 4/4 [00:03<00:00, 1.33it/s]
```

```
Valid loss : 1.0678082842093248
```

```
Valid dice score : 0.1276569962501526
```

```
100%|██████████| 26/26 [00:14<00:00, 1.83it/s]
```

```
Train loss : 0.8607734304208022
```

```
Train dice score : 0.02353499084711075
```

```
100%|██████████| 4/4 [00:03<00:00, 1.31it/s]
```

```
Valid loss : 43.77381962308517
```

```
Valid dice score : 0.0004553273902274668
```

```
100%|██████████| 26/26 [00:14<00:00, 1.83it/s]
```

```
Train loss : 0.8768505454063416
```

```
Train dice score : 0.02549426071345806
```

```
100%|██████████| 4/4 [00:03<00:00, 1.30it/s]
```

```
Valid loss : 6.183049529790878
```

```
Valid dice score : 0.0
```

```
100%|██████████| 26/26 [00:14<00:00, 1.80it/s]
```

```
Train loss : 0.8849088240128297
```

```
Train dice score : 0.022647453472018242
```

```
100%|██████████| 4/4 [00:03<00:00, 1.31it/s]
```

```
Valid loss : 1.0279537324721997
```

```
Valid dice score : 0.011607266962528229
```

```
100%|██████████| 26/26 [00:14<00:00, 1.82it/s]
```

```

Train loss : 0.8407229781150818
Train dice score : 0.011387759819626808
100%|██████████| 4/4 [00:03<00:00, 1.27it/s]
Valid loss : 255.50482940673828
Valid dice score : 0.0
100%|██████████| 26/26 [00:14<00:00, 1.81it/s]
Train loss : 0.7958291104206672
Train dice score : 0.0036396444775164127
100%|██████████| 4/4 [00:03<00:00, 1.33it/s]
Valid loss : 1991.4961327772874
Valid dice score : 0.0
100%|██████████| 26/26 [00:14<00:00, 1.79it/s]
Train loss : 0.8863115333593808
Train dice score : 0.00577913923189044
100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
Valid loss : 42.11325320830712
Valid dice score : 0.12204891443252563

```

```
Out[ ]: (0.001, 0.001)
```

```
In [ ]: best_lr, best_weight_decay
```

```
Out[ ]: (0.001, 0.001)
```

Training

```
In [ ]: LEARNING_RATE = best_lr #1e-3
WEIGHT_DECAY = best_weight_decay #5e-4
```

```
In [ ]: LEARNING_RATE = 1e-3
WEIGHT_DECAY = 5e-4
```

```
In [ ]: model = UNet(out_channels=5+1)
model.to(DEVICE)

optimizer = torch.optim.AdamW(params=model.parameters(),
                               lr=LEARNING_RATE,
                               weight_decay=WEIGHT_DECAY)

criterion = nn.CrossEntropyLoss()
```

```
In [ ]: NUM_EPOCHS = 30
```

```
In [ ]: best_score = -np.inf

train_loss_list = []
train_score_list = []

valid_loss_list = []
valid_score_list = []
```

```

for epoch in range(NUM_EPOCHS):

    print("Epoch = ", epoch)

    train_loss, train_score = train_epoch(
        model,
        optimizer=optimizer,
        criterion=criterion,
        data_loader=train_loader,
        device=DEVICE
    )

    valid_loss, valid_score = valid_epoch(
        model,
        optimizer=optimizer,
        criterion=criterion,
        data_loader=valid_loader,
        device=DEVICE
    )

    train_loss_list.append(train_loss)
    train_score_list.append(train_score)

    valid_loss_list.append(valid_loss)
    valid_score_list.append(valid_score)

    if valid_score > best_score:
        best_score = valid_score
        if SAVE_MODEL:
            print("Model improved, saving...")
            checkpoint = {
                "state_dict": model.state_dict(),
                "optimizer": optimizer.state_dict(),
            }
            save_checkpoint(checkpoint, filename=f"/content/drive/MyDrive/Unet/MINI")
        print('\n')

```

Epoch = 0

100%|██████████| 26/26 [00:14<00:00, 1.76it/s]

Train loss : 1.2476573013342345

Train dice score : 0.08314365893602371

100%|██████████| 4/4 [00:03<00:00, 1.27it/s]

Valid loss : 1.167629361152649

Valid dice score : 0.034658610820770264

Model improved, saving...

=> Saving checkpoint

Epoch = 1

100%|██████████| 26/26 [00:14<00:00, 1.80it/s]

Train loss : 0.8419976830482483

Train dice score : 0.039488714188337326

100%|██████████| 4/4 [00:03<00:00, 1.26it/s]

Valid loss : 0.9681942600470322
Valid dice score : 0.10438269376754761
Model improved, saving...
=> Saving checkpoint

Epoch = 2

100%|██████████| 26/26 [00:15<00:00, 1.71it/s]

Train loss : 0.7296106746563544

Train dice score : 0.03503970056772232

100%|██████████| 4/4 [00:03<00:00, 1.21it/s]

Valid loss : 0.8928506511908311

Valid dice score : 0.02458646148443222

Epoch = 3

100%|██████████| 26/26 [00:14<00:00, 1.78it/s]

Train loss : 0.6906019552395894

Train dice score : 0.027427153661847115

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.8520097549145038

Valid dice score : 0.017835283651947975

Epoch = 4

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]

Train loss : 0.6704001931043772

Train dice score : 0.03240266442298889

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.8206080496311188

Valid dice score : 0.008662436157464981

Epoch = 5

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]

Train loss : 0.6554523809598043

Train dice score : 0.015467824414372444

100%|██████████| 4/4 [00:03<00:00, 1.30it/s]

Valid loss : 0.8272094749487363

Valid dice score : 0.004202330484986305

Epoch = 6

100%|██████████| 26/26 [00:14<00:00, 1.78it/s]

Train loss : 0.6437448389255084

Train dice score : 0.03625857084989548

100%|██████████| 4/4 [00:03<00:00, 1.31it/s]

Valid loss : 0.7809476004197047

Valid dice score : 0.01859412156045437

Epoch = 7

100%|██████████| 26/26 [00:14<00:00, 1.75it/s]

Train loss : 0.6243345531133505
Train dice score : 0.049412090331315994

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
Valid loss : 0.8119312799893893
Valid dice score : 0.07507951557636261

Epoch = 8

100%|██████████| 26/26 [00:14<00:00, 1.77it/s]
Train loss : 0.6345491099816102
Train dice score : 0.041943199932575226

100%|██████████| 4/4 [00:03<00:00, 1.30it/s]
Valid loss : 0.7943372909839337
Valid dice score : 0.009358283132314682

Epoch = 9

100%|██████████| 26/26 [00:14<00:00, 1.80it/s]
Train loss : 0.6058349357201502
Train dice score : 0.05595178157091141

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
Valid loss : 0.9266647776732078
Valid dice score : 0.011807403527200222

Epoch = 10

100%|██████████| 26/26 [00:14<00:00, 1.86it/s]
Train loss : 0.6186931660542121
Train dice score : 0.06704144924879074

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
Valid loss : 0.8606838354697595
Valid dice score : 0.16748584806919098
Model improved, saving...
=> Saving checkpoint

Epoch = 11

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]
Train loss : 0.5735325893530479
Train dice score : 0.10405883938074112

100%|██████████| 4/4 [00:03<00:00, 1.27it/s]
Valid loss : 0.6929159645850842
Valid dice score : 0.17510759830474854
Model improved, saving...
=> Saving checkpoint

Epoch = 12

100%|██████████| 26/26 [00:14<00:00, 1.81it/s]
Train loss : 0.5868111241322297
Train dice score : 0.10840272158384323

100%|██████████| 4/4 [00:03<00:00, 1.26it/s]

Valid loss : 0.8274962833294501
Valid dice score : 0.0213769618421793

Epoch = 13

100%|██████████| 26/26 [00:14<00:00, 1.78it/s]

Train loss : 0.5971231345946972

Train dice score : 0.06372759491205215

100%|██████████| 4/4 [00:03<00:00, 1.30it/s]

Valid loss : 0.7143219411373138

Valid dice score : 0.08615143597126007

Epoch = 14

100%|██████████| 26/26 [00:14<00:00, 1.77it/s]

Train loss : 0.5770060878533584

Train dice score : 0.10685397684574127

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.6812461866782262

Valid dice score : 0.1448332518339157

Epoch = 15

100%|██████████| 26/26 [00:14<00:00, 1.85it/s]

Train loss : 0.5746508481410834

Train dice score : 0.10862919688224792

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.7197995552649865

Valid dice score : 0.042692072689533234

Epoch = 16

100%|██████████| 26/26 [00:14<00:00, 1.80it/s]

Train loss : 0.5725673081783148

Train dice score : 0.09520527720451355

100%|██████████| 4/4 [00:03<00:00, 1.30it/s]

Valid loss : 0.757398050564986

Valid dice score : 0.11068412661552429

Epoch = 17

100%|██████████| 26/26 [00:15<00:00, 1.72it/s]

Train loss : 0.5729008809878275

Train dice score : 0.07995118200778961

100%|██████████| 4/4 [00:03<00:00, 1.30it/s]

Valid loss : 0.7452404040556687

Valid dice score : 0.0489567294716835

Epoch = 18

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]

Train loss : 0.5641662730620458

Train dice score : 0.1156081035733223

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.6806393414735794

Valid dice score : 0.17179416120052338

Epoch = 19

100%|██████████| 26/26 [00:14<00:00, 1.81it/s]

Train loss : 0.5555833142537338

Train dice score : 0.10239347815513611

100%|██████████| 4/4 [00:03<00:00, 1.26it/s]

Valid loss : 0.89966231240676

Valid dice score : 0.00428706593811512

Epoch = 20

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]

Train loss : 0.5545899684612567

Train dice score : 0.10585465282201767

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.691686905347384

Valid dice score : 0.103823222219944

Epoch = 21

100%|██████████| 26/26 [00:14<00:00, 1.76it/s]

Train loss : 0.5474112675740168

Train dice score : 0.11742405593395233

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]

Valid loss : 0.671754682293305

Valid dice score : 0.1327986717224121

Epoch = 22

100%|██████████| 26/26 [00:14<00:00, 1.79it/s]

Train loss : 0.5477876078623992

Train dice score : 0.11006250232458115

100%|██████████| 4/4 [00:03<00:00, 1.29it/s]

Valid loss : 0.6668556928634644

Valid dice score : 0.11167226731777191

Epoch = 23

100%|██████████| 26/26 [00:14<00:00, 1.84it/s]

Train loss : 0.5439692930533335

Train dice score : 0.12616916000843048

100%|██████████| 4/4 [00:02<00:00, 1.34it/s]

Valid loss : 0.7802888051821635

Valid dice score : 0.009086656384170055

Epoch = 24

100%|██████████| 26/26 [00:14<00:00, 1.83it/s]

Train loss : 0.53352002455638
Train dice score : 0.1252955198287964
100%|██████████| 4/4 [00:02<00:00, 1.33it/s]
Valid loss : 0.6909337238623545
Valid dice score : 0.11758770048618317

Epoch = 25

100%|██████████| 26/26 [00:14<00:00, 1.82it/s]
Train loss : 0.5665271706306018
Train dice score : 0.11716224253177643
100%|██████████| 4/4 [00:03<00:00, 1.31it/s]
Valid loss : 0.7743223034418546
Valid dice score : 0.09266463667154312

Epoch = 26

100%|██████████| 26/26 [00:14<00:00, 1.80it/s]
Train loss : 0.5771467009415994
Train dice score : 0.10959058254957199
100%|██████████| 4/4 [00:03<00:00, 1.33it/s]
Valid loss : 0.7305118796917108
Valid dice score : 0.06824445724487305

Epoch = 27

100%|██████████| 26/26 [00:14<00:00, 1.84it/s]
Train loss : 0.5542146724004012
Train dice score : 0.10607340186834335
100%|██████████| 4/4 [00:03<00:00, 1.31it/s]
Valid loss : 0.6294951117955722
Valid dice score : 0.17580631375312805
Model improved, saving...
=> Saving checkpoint

Epoch = 28


100%|██████████| 26/26 [00:14<00:00, 1.76it/s]
Train loss : 0.548497314636524
Train dice score : 0.12256655097007751
100%|██████████| 4/4 [00:03<00:00, 1.25it/s]
Valid loss : 0.7040876728984026
Valid dice score : 0.05634910613298416

Epoch = 29


100%|██████████| 26/26 [00:14<00:00, 1.83it/s]
Train loss : 0.5379140629218175
Train dice score : 0.11529923230409622
100%|██████████| 4/4 [00:02<00:00, 1.33it/s]

Valid loss : 0.6927718325303152
Valid dice score : 0.05446914583444595

```
In [ ]: plt.plot(train_loss_list, label='train_loss', marker='*')
plt.plot(valid_loss_list, label='val_loss', marker='*')
plt.title('Loss per epoch'); plt.ylabel('loss');
plt.xlabel('epoch')
plt.legend(), plt.grid()
plt.show()
```



```
In [ ]: plt.plot(train_score_list, label='train_score', marker='*')
plt.plot(valid_score_list, label='val_score', marker='*')
plt.title('score per epoch'); plt.ylabel('score');
plt.xlabel('epoch')
plt.legend(), plt.grid()
plt.show()
```



Test

```
In [ ]: best_check_point = torch.load("/content/drive/MyDrive/Unet/MINI_Unet_Best.pth.tar")
```

```
In [ ]: model.load_state_dict(best_check_point['state_dict'])  
model.eval()
```

```

Out[ ]: UNet(
  (pool): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil
_mode=False)
  (conv1): encoding_block(
    (conv): Sequential(
      (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=F
alse)
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_sta
ts=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=
False)
      (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_sta
ts=True)
      (5): ReLU(inplace=True)
    )
  )
  (conv2): encoding_block(
    (conv): Sequential(
      (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
      (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
      (5): ReLU(inplace=True)
    )
  )
  (conv3): encoding_block(
    (conv): Sequential(
      (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
      (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
      (5): ReLU(inplace=True)
    )
  )
  (conv4): encoding_block(
    (conv): Sequential(
      (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bia
s=False)
      (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)

```

```

        (5): ReLU(inplace=True)
    )
)
(conv5): encoding_block(
  (conv): Sequential(
    (0): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
  )
)
(conv6): encoding_block(
  (conv): Sequential(
    (0): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
  )
)
(conv7): encoding_block(
  (conv): Sequential(
    (0): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
  )
)
(conv8): encoding_block(
  (conv): Sequential(
    (0): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```



```

        (5): ReLU(inplace=True)
    )
)
(tconv1): ConvTranspose2d(1024, 512, kernel_size=(2, 2), stride=(2, 2))
(tconv2): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
(tconv3): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))
(tconv4): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))
(bottleneck): encoding_block(
  (conv): Sequential(
    (0): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
  )
)
(final_layer): Conv2d(64, 6, kernel_size=(1, 1), stride=(1, 1))
)

```

```

In [ ]: test_loss, test_score = valid_epoch(
        model,
        optimizer=optimizer,
        criterion=criterion,
        data_loader=test_loader,
        device=DEVICE
    )

```

```

100%|██████████| 4/4 [00:03<00:00, 1.32it/s]
Valid loss : 0.6054971676606399
Valid dice score : 0.20220284163951874

```

```

In [ ]: print("Test Loss : ", test_loss)
        print("Test Score : ", test_score)

```

```

Test Loss : 0.6054971676606399
Test Score : 0.20220284163951874

```

```

In [ ]: softmax = nn.Softmax(dim=1)

```

```

In [ ]: # Disable gradient calculation to speed up inference
        with torch.no_grad():
            # Loop over the data loader
            for batch_idx, (inputs, targets) in enumerate(tqdm(test_loader)):
                # Move the inputs and targets to the device
                inputs, targets = inputs.to(DEVICE), targets.to(DEVICE)

                # Forward pass
                outputs = model(inputs)
                outputs = torch.argmax(softmax(outputs), axis=1)

                for i in range(inputs.shape[0]):

```

```

image = inputs[i].to('cpu')
mask = targets[i].to('cpu')
pred_mask = outputs[i].to('cpu')

fig , ax = plt.subplots(1, 3, figsize=(18, 18))
softmax = nn.Softmax(dim=1)

image = image.permute(1, 2, 0).numpy()
mask = mask.numpy()
pred_mask = pred_mask.numpy()

ax[0].imshow(image)
ax[0].set_title("Image")

ax[1].imshow(mask)
ax[1].set_title("Mask")


ax[2].imshow(pred_mask)
ax[2].set_title("Predicted Mask")

plt.show()


# break

```


0%| | 0/4 [00:00<?, ?it/s]WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




25%|██████████| 1/4 [00:03<00:10, 3.36s/it]WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




50%|██████████| 2/4 [00:07<00:07, 3.58s/it]WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



75%|██████████| 3/4 [00:10<00:03, 3.49s/it]WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



100%|██████████| 4/4 [00:11<00:00, 2.83s/it]

